

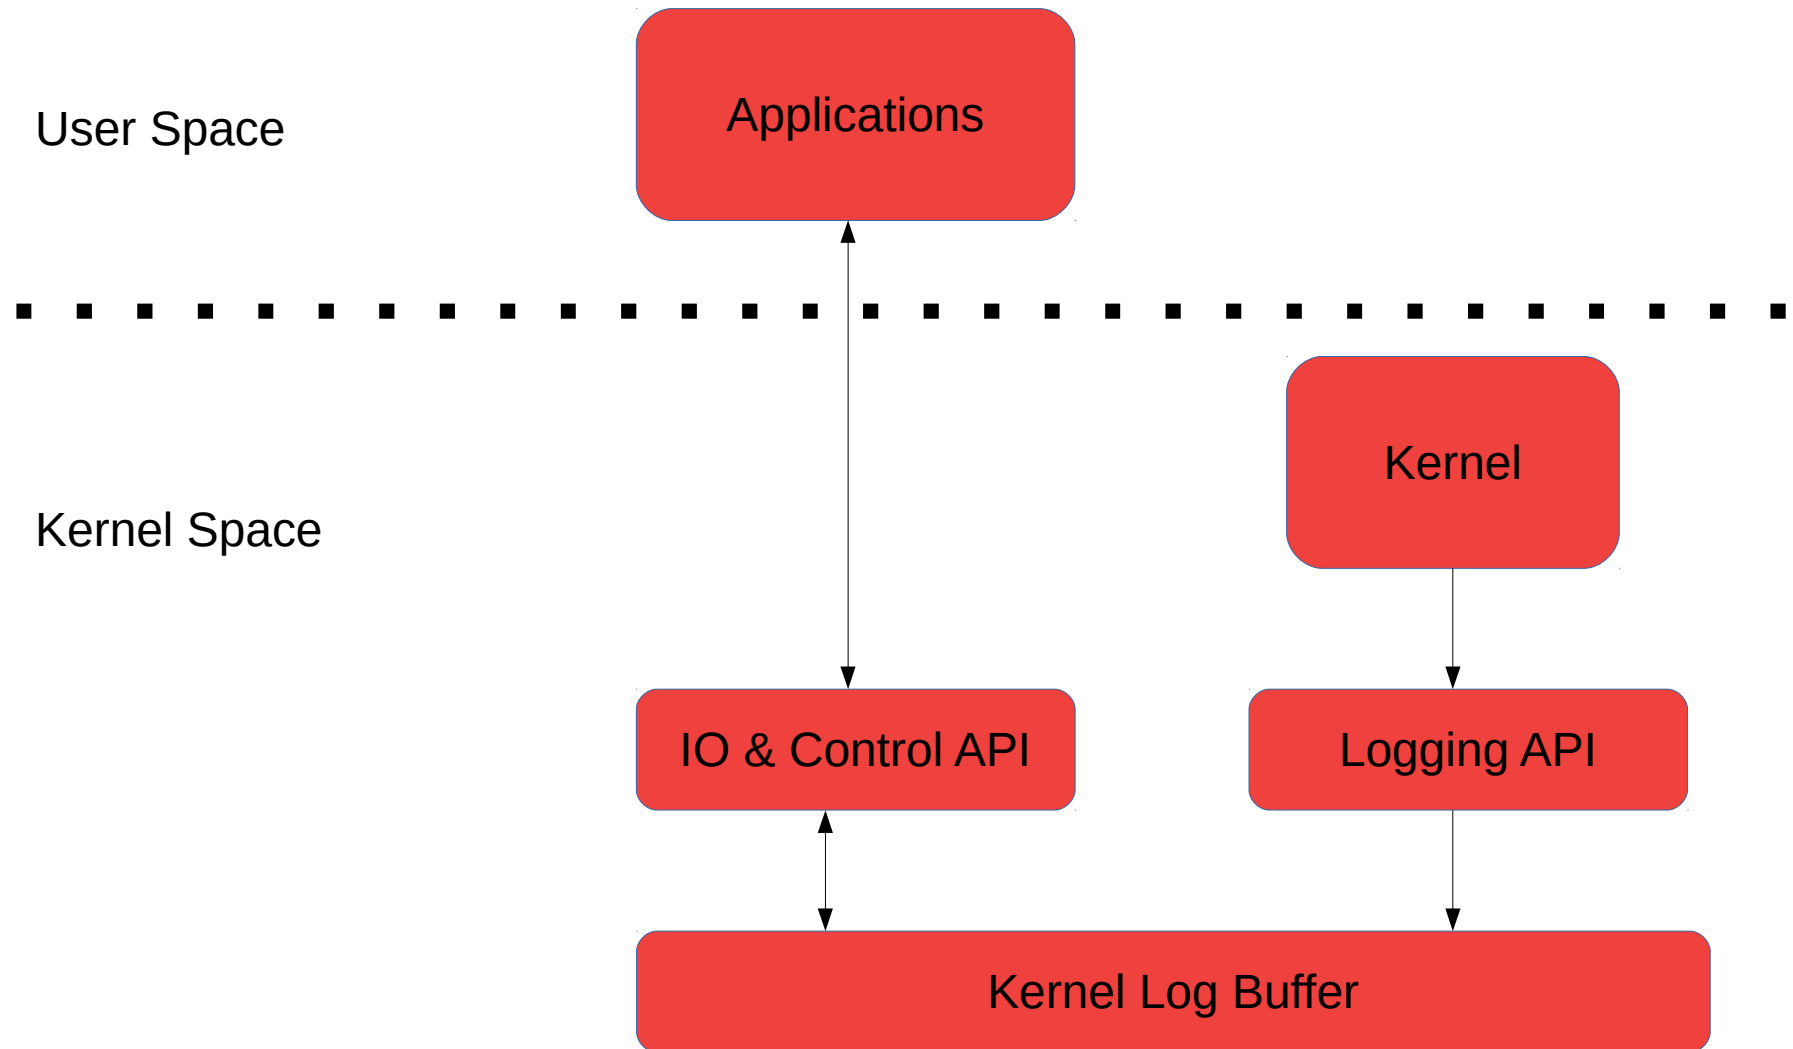
Kernel Debugging



What to Expect

- Debugging by printing
- DebugFS
- Kgdb
- Analyzing the oops
- Analyzing the boot up time
- Other debugging mechanisms

Kernel Logging System Architecture



Kernel Log buffer

- Default size is 64KB
- For modifying the size
 - Kernel Config Option
 - `CONFIG_LOG_BUF_SHIFT=n`
 - Menuconfig -> General Setup
 - Uboot bootargs: `log_buf_len=n`
 - Buffer Size = 2^n
 - $n = 16$: 64 KB
 - $n = 17$: 128 KB

Debugging by printing

- Simplest & the most commonly used debug method
 - `printk(KERN_ALERT "reached line %d in function %s\n", __LINE__, __func__);`
- Limitations
 - Recompile & reboot every time, the new section needs debugging
 - Prints are relatively resource intensive
 - Might not help if the bug is related to timing/resource contention

Debugging by printing ...

- Using `pr_*` family of functions
 - Shorthand definitions for the respective `printk` call
 - `pr_emerg`, `pr_alert`, `pr_crit`, `pr_err`, `pr_debug`, `pr_devel`
 - `pr_devel` and `pr_debug` are replaced with `printk(KERN_DEBUG ..)` if the kernel was compiled with `DEBUG`, otherwise replaced with empty statement
- Example
 - `pr_emerg("Error in allocation\n");`

Debug prints

- `dev_*`
 - Special version of `printk` wrapper routines for device drivers
 - Show the extra information
 - `dev_emerg`, `dev_crit`, `dev_alert`, `dev_err`, `dev_warn`, `dev_notice`
 - Used when printing something related to devices
- `dev_dbg` & `pr_debug`
 - Not compiled by default
- `#define DEBUG` at the beginning of the driver
- Using `ccflags-{CONFIG_DRIVER} += -DDEBUG` in the Makefile

Log Level

- Log level in the message defaults to `DEFAULT_MESSAGE_LOGLEVEL`
 - Can be set via the `CONFIG_DEFAULT_MESSAGE_LOGLEVEL` kernel config option (make menuconfig-> Kernel Hacking -> Default message log level)
 - The log level is used by the kernel to determine the importance of a message and to decide whether it should be presented to the user immediately
 - To determine your current `console_loglevel` you simply enter:
 - `cat /proc/sys/kernel/printk`

7	4	1	7
current	default	minimum	boot-time-default
 - `echo 8 > /proc/sys/kernel/printk`
 - `#set console_loglevel` to print `KERN_WARNING` (4) or more severe messages
 - `# dmesg -n 5`

Rate Limiting & one time messages

- Inserting a `printk` in a section which gets called quite often might result in a severe performance impact
 - Could overwrite & spam the kernel buffer
- `printk_once(...)`
 - no matter how often you call it, it prints once and never again
- `#include <kernel/ratelimit.h>`
- `printk_ratelimited(...)`
 - it prints by default not more than 10 times in every 5 seconds (for each function it is called in).

Printing from user space

- To annotate, its quite useful to insert some messages in the kernel log buffer
 - echo "Hello Kernel-World" > /dev/kmsg
 - Prints with the default log level
 - echo "<2>Writing critical printk messages from userspace" > /dev/kmsg
 - To issue a KERN_CRIT message
 - Example:
 - echo "#### TESTNOTE: unplugged thumb drive" > /dev/kmsg
 - echo "#### TESTNOTE: waited for a couple seconds" > /dev/kmsg
 - echo "#### TESTNOTE: re-plugged thumb drive" > /dev/kmsg

Printing buffers as hex

- `print_hex_dump_bytes(const char *prefix_str, int prefix_type, const void *buf, size_t len)`
- `static inline void print_hex_dump(const char *level, const char *prefix_str, int prefix_type, int rowsize, int groupsize, const void *buf, size_t len, bool ascii)`
- prints a buffer as hex values to the kernel log buffer (with level `KERN_DEBUG`)
- Useful for creating the memory dumps
- Example
 - `char mybuf[] = "abcdef";`
 - `print_hex_dump_bytes("", DUMP_PREFIX_NONE, mybuf, ARRAY_SIZE(mybuf));`
 - dmesg output:
 - 61 62 63 64 65 66 00 abcdef.

Dynamic debug

- Can be used to enable/disable debug information dynamically
 - Kernel needs to be compiled with `CONFIG_DYNAMIC_DEBUG`
 - Useful tool to only get the debug messages you are interested in
 - `pr_debug()/dev_dbg()` and `print_hex_dump_debug()/print_hex_dump_bytes()` calls can be dynamically enabled per-callsite
 - If `CONFIG_DYNAMIC_DEBUG` is not set, `print_hex_dump_debug()` is just shortcut for `print_hex_dump(KERN_DEBUG)`.

Dynamic Debug ...

- Simple query language allows turning on and off debugging statements by matching any combination of:
 - source filename
 - function name
 - Line number (including ranges of line numbers)
 - module name
- Enable debug messages during boot process
 - `dyndbg="QUERY" <--` for kernel
 - `module.dyndbg="QUERY" < --` for module

Dynamic debug control options

- Using DebugFS
 - `mount -t debugfs none /sys/kernel/debug/`
 - `# cd /sys/kernel/debug/dynamic_debug/`
 - `# echo "file xxx.c +p" > control`
 - `# echo "file svcsock.c line 1603 +p" > control`
 - `# echo "file drivers/usb/core/* +p" > control`
 - `# echo "file xxx.c -p" > control`
- `uboot bootargs`
 - `dyndbg="QUERY" <-- for kernel`
 - `module.dyndbg="QUERY" < -- for module`

Using DebugFS

- A simple memory based filesystem designed specifically to debug Linux kernel code
- Helps kernel developers export large amount of debug data into user space
- Kernel Configuration: CONFIG_DEBUG_FS
 - Kernel hacking -> Debug Filesystem
- Mount debugfs with command
 - `mount -t debugfs nodev /sys/kernel/debug`

DebugFS API

- Create a subdirectory in /sys/kernel/debug
 - Struct dentry *debugfs_create_dir(const char *name, struct dentry *parent)
 - Expose an integer using file in DebugFS
 - Struct dentry *debugfs_create_u8(const char *name, mode_t mode, struct dentry *parent, u8 *value)
 - Expose a binary blob
 - Struct dentry *debugfs_create_blob(const char *name, mode_t mode, struct dentry *parent, struct debugfs_blob_wrapper *blob)

Kernel Probes

- Mechanism to write the modules that can add debug information to the kernel
- An alternative to building custom kernels or custom modules
- Dynamically breaks into any kernel routine and can collect debugging and performance information non-disruptively.
- Typical use case
 - Debugging a remote machine where dmesg is not enough to debug. Build a kprobe module & then insmod on remote machine
- Types
 - jprobes
 - Function is called on the entry to the routine. All the arguments to the routine are passed
 - kprobes
 - Any arbitrary kernel instruction can be probed. A function is called passing the registers
 - Kretprobe
 - Call a function on the exit from the routine. The registers are passed

ftrace

- Stands for Function Tracer
- Can be used for
 - Debugging Linux Kernel
 - Analyzing latencies in Linux Kernel
 - Learn & observe the flow of Linux Kernel
 - Trace Context switches
 - Length of the time the interrupts are disabled
- Kernel Configuration
 - CONFIG_FTRACE --> "Tracers"
 - CONFIG_FUNCTION_TRACER --> Kernel Function Tracer
 - CONFIG_FUNCTION_GRAPH_TRACER --> Kernel Function Graph Tracer
 - CONFIG_DYNAMIC_TRACE --> Enable/Disable ftrace dynamically

ftrace operations

- Mount tracefs
 - Adding the entry into the fstab
 - `tracefs /sys/kernel/tracing tracefs_defaults 0 0`
 - Using the mount command
 - `mount -t tracefs nodev /sys/kernel/tracing`
- `available_tracers`
 - Lists what all tracers have been enabled in the kernel configuration
- `current_tracer`
 - The tracer currently is running
- `trace`
 - Contains the tracing data in human readable format
- `tracing_on`
 - Enable/disable writing tracing data to ring buffer (ftrace uses a separate ring buffer to store tracing data)
- To enable function tracer
 - `echo "function" > current_tracer`

Function Graph

- Is Used to
 - track the entry of the function
 - track the exit of the function
 - find the Execution Time
 - get the CPU on which it is running
- Useful for following the flow of execution within the kernel

Tracing a specific process

- Steps to trace the process
 - Disable tracing
 - `echo "nop" > current_tracer`
 - Echo pid of the process which you want to trace in "set_ftrace_pid" file
 - `echo "2588" > set_ftrace_pid`
 - Enable the function tracer
 - `echo "function_graph" > current_tracer`

Dynamic Tracing

- Used to filter just the function we need and eliminate those we don't need
- Can be done with the file 'set_ftrace_filter'
 - cat available_filter_functions
 - echo vmalloc_* > set_ftrace_filter
- <https://01.org/linuxgraphics/gfx-docs/drm/trace/ftrace.html>

MMIO tracing

- Refer
<https://www.kernel.org/doc/Documentation/trace/mmiotrace.txt>

trace_printk

- Limitations with printk
 - Using printk in interrupt context can create a live lock
 - The bug might disappear if printk is added, in case, its time sensitive
 - May take several milliseconds when writing to the console
- trace_printk advantages
 - Writing will be in the order of microseconds as it writes to a ring buffer instead of console
 - Can be used in any context (interrupt, scheduler, NMI Code)
 - Can be read via the 'trace' file

perf Tool

- A profiling tool which offers support for tracing applications and also inspecting the general aspects of the system
- Allows to take a look at what functions are being called at a given point
- Allows us to take a peak at where the kernel is spending most of the time, prints out the call stack and in general logs what the cpu is running
- `sudo perf record -a -g`
- `perf report --header -F overhead,comm,parent`
- `sudo perf timechart record`
- `sudo perf timechart -->` Generates the .svg file

Kernel Debuggers

- Two debugger front ends – KDB & KGDB
- KDB
 - Simplistic shell-style interface
 - Used to inspect memory, registers, process lists, dmesg and even set breakpoints to stop in a certain location
 - Not a source level debugger
 - Aimed at doing some analysis for developing or diagnosing kernel problems
- KGDB
 - To used as a source level debugger
 - Used along with the gdb to debug a linux kernel
 - Gdb can be used to break-in to the kernel to inspect memory, variables and look through call stack information

Kernel GDB

- Provides an interface to gdb via its remote serial protocol
- Implements a gdb stub that communicates to the cross gdb running on host
- Kernel Configuration
 - CONFIG_FRAME_POINTER=y
 - CONFIG_KGDB=y
 - CONFIG_KGDB_SERIAL_CONSOLE

KGDB setup

- Configure KGDB from command line
 - `kgdboc = <tty-device>, <bauds>`
 - Add `kgdbwait` to make `kgdb` wait for the debugger connection
- On the host
 - `arm-linux-gdb ./vmlinux`
 - Set `remotebaud 115200`
 - Target remote `/dev/ttyUSB0`

Kernel Oops

- An exception in the kernel code
- Kernel dumps this message when it finds something faulty
- Contains the processor state & the CPU registers of when the fault occurred
- The offending process gets killed without even releasing the locks or cleaning up the data structures
- System cannot be trusted further, once the oops have happened

Analyzing Kernel oops

- BUG: What caused the oops
- PC: Instruction pointer
- Internal error: [#1] SMP – This is error code in hex
 - Varies as per architecture
- CPU 1 – the CPU on which the error occurred
- Call Trace - the list of functions being called just before the Oops occurred
- Code: The Code is a hex-dump of the section of machine code that was being run at the time the Oops occurred

Debugging an Oops dump

- `gdb test.ko`
- `(gdb) add-symbol-file test.o <address>`
 - Add the symbol file to the debugger
 - The address of the test section of the module
 - `cat /sys/module/test/sections/.init.text`
- `(gdb) disassemble my_oops_init`
 - `my_oops_init` is the offending function
 - We can get it from the PC
- Add the starting address & the offset to pin point the actual line of offending code
- `(gdb) list *(address)`

Analyzing the boot up time

- Variety of tools available to measure the boot up time for the linux system
- grabserial
 - One of the simplest tool
 - Reads the serial port and writes the data to the standard output
 - `grabserial -d /dev/ttyUSB0 -t`

Other debugging mechanisms

- Adding the ioctl commands for debugging mechanisms
- Adding entries in the proc filesystem
- Adding debugging entries in sys filesystem

What all did we learn?

- Debugging by printing
- DebugFS
- Kgdb
- Analyzing the oops
- Other Debugging Mechanisms